

Random For(r)est

Sebastian Heilmann

19 8 2020

1 Explore data

In this script, we will use linguistic features to predict the speaker (Forrest Gump vs anyone else) of spoken lines in the movie “Forrest Gump” by Robert Zemeckis.

Let’s start by exploring some sentences.

```
as.tbl(head(forrest))
```

```
## # A tibble: 6 x 32
##   run  onset duration offset time_to_next overlapping person FG   kind
##   <chr> <dbl>    <dbl> <dbl> <chr>          <chr>    <chr> <chr> <chr>
## 1 run-1 1207      1.4    2607 1.96          False  ERZAE... 1  spee...
## 2 run-1 3167      3.84   7007 4.16          False  ERZAE... 1  spee...
## 3 run-1 7327      1.6    8927 2.08          False  ERZAE... 1  spee...
## 4 run-1 9407      0.68  10087 6.6          False  ERZAE... 1  spee...
## 5 run-1 16007     3.48  19487 4.76          False  ERZAE... 1  spee...
## 6 run-1 20767     145.  165567 2.08          False  OST      1  song
## # ... with 23 more variables: prosody <chr>, sentence_grammar <chr>, text <chr>,
## #   comp_score <dbl>, word_count <int>, lexfreq_norm_log_sum <dbl>,
## #   lexfreq_norm_log_mean <dbl>, lexmin <dbl>, lexmax <dbl>, sent_sum <dbl>,
## #   isforrest <chr>, satzid <fct>, open <int>, closed <int>, undefined <int>,
## #   document <chr>, CTTR <dbl>, n_lowers <dbl>, n_caps <dbl>, n_uq_words <dbl>,
## #   n_uq_chars <dbl>, n_charsperword <dbl>, depth <dbl>
```

```
set.seed(51)
forrest %>%
  filter(!is.na(isforrest)) %>%
  group_by(isforrest) %>%
  sample_n(4) %>%
  select(isforrest, text)
```

```
## # A tibble: 8 x 2
## # Groups:   isforrest [2]
##   isforrest text
##   <chr>    <chr>
## 1 FORREST  Wir waren wie eine Familie
## 2 FORREST  Und das Abendessen
## 3 FORREST  Wir waren wie Pech und Schwefel
## 4 FORREST  Sie wollte dass ich die beste Schulbildung bekomme
## 5 OTHER    Es ist hell
## 6 OTHER    Everybodys Talkin
## 7 OTHER    Denk immer an das was ich dir gesagt habe
## 8 OTHER    Sie trägt noch ihr türkisfarbenes Kellnerinnen-Kleid
```

We have a lot of features! But let’s also try to take the sentence length into account and recalculate / add some features.

```
forrest <- forrest %>%
  mutate(sentiment = sent_sum / word_count,
         nodes = comp_score / word_count,
         open_words = open / word_count,
         closed_words = closed / word_count,
         speech_rate = word_count / duration)
```

Starting with exploratory data analysis is always important before modeling. Let’s therefore make one nice plot to explore the relationships in this data. Are there any text features that distinguish Forrest Gump from other speakers?

Overview of predictors

Mean text predictors per sentence



Data: Master Thesis

We can see differences in lexical features and character-features especially. We will now create a separate data set with those variables.

```
# select relevant variables
forrest_clean <- forrest_clean %>%
  select(c(CTTR:speech_rate, duration, word_count, lexfreq_norm_log_mean:lexmax,isforrest)) %>%
  filter(!is.infinite(speech_rate)) %>%
  na.omit()
```

2 Build a random forest model

2.1 Split data into training and test data set

We will start by splitting our data into a training and testing set. Then we will create cross-validation resamples of the training data to evaluate our models.

```
set.seed(123)
forrest_split <- initial_split(forrest_clean, strata = isforrest)
forrest_train <- training(forrest_split)
forrest_test <- testing(forrest_split)

set.seed(123)
forrest_folds <- vfold_cv(forrest_train, strata = isforrest)
```

2.2 Define a tidmodel recipe and workflow.

We now want to tune some hyperparameters. Because we use a random forest model, we don't need to do much preprocessing. That means, that we don't need to

worry about centering or scaling our data.

```
forrest_recipe <-  
  recipe(isforrest ~ ., data = forrest_clean) %>%  
  step_naomit(all_predictors()) %>%  
  step_zv(all_predictors())  
  
rf_model <-  
  # specify that the model is a random forest  
  rand_forest() %>%  
  # specify parameters that we want to tune  
  set_args(mtry = tune(), trees = tune(), min_n = tune()) %>%  
  # select the engine that underlies the model  
  set_engine("ranger", importance = "impurity") %>%  
  # choose binary classification mode  
  set_mode("classification")  
rf_model
```

```
## Random Forest Model Specification (classification)  
##  
## Main Arguments:  
##   mtry = tune()  
##   trees = tune()  
##   min_n = tune()  
##  
## Engine-Specific Arguments:  
##   importance = impurity  
##  
## Computational engine: ranger
```

```
rf_workflow <- workflow() %>%  
  # add the recipe  
  add_recipe(forrest_recipe) %>%  
  # add the model  
  add_model(rf_model)  
  
rf_workflow
```

```
## == Workflow ==  
## Preprocessor: Recipe  
## Model: rand_forest()  
##  
## — Preprocessor —  
## 2 Recipe Steps  
##  
## ● step_naomit()  
## ● step_zv()  
##  
## — Model —  
## Random Forest Model Specification (classification)  
##  
## Main Arguments:  
##   mtry = tune()  
##   trees = tune()  
##   min_n = tune()  
##  
## Engine-Specific Arguments:  
##   importance = impurity  
##  
## Computational engine: ranger
```

2.3 Tune hyperparameters


We will use `tune_grid()` with our tuneable workflow (and our grid of parameters (`mtry`, `trees`, `min_n`) and our resamples to try. Most importantly, but optional, we will call `doParallel::registerDoParallel()` to parallelize R code on our machine.

```
doParallel::registerDoParallel()  
  
set.seed(234)  
tune_res <- tune_grid(  
  rf_workflow,  
  resamples = forrest_folds,  
  grid = 20  
)
```

Let's try a visualization to understand our results.

Overview of hyperparameters

It looks like lower values of mtry (number of predictors sampled for splitting at each node) and higher values of trees perform well.



What are the best performing sets of parameters?

```
show_best(tune_res, "roc_auc")
```

```
## # A tibble: 5 x 9
##   mtry trees min_n .metric .estimator mean      n std_err .config
##   <int> <int> <int> <chr>   <chr>   <dbl> <int>   <dbl> <chr>
## 1     2  1236     3 roc_auc binary  0.800    10 0.00649 Model12
## 2     5   940    15 roc_auc binary  0.798    10 0.00644 Model11
## 3     4   310    16 roc_auc binary  0.798    10 0.00620 Model103
## 4     7  1605    40 roc_auc binary  0.797    10 0.00623 Model101
## 5     6   478    20 roc_auc binary  0.796    10 0.00648 Model105
```

We will tune our hyperparameters once more, but in a more fine grained way.

```
set.seed(123)
# set tuning parameters to try
rf_grid <- expand_grid(mtry = c(2,3,4), trees = c(1300, 1500, 2000), min_n = c(4,6,8,10,20,30,32,34,36))
# apply model and extract results
rf_tune_results <- rf_workflow %>%
  tune_grid(resamples = forrest_folds,
    grid = rf_grid,
    # choose some metrics
    metrics = metric_set(accuracy, roc_auc, sens, spec)
  )
```

Let's have a look at all combinations of parameters.

Random Forest

ROC-AUC Scores of random forest models with varying number of trees, split-nodes and terminal nodes



3 Finalize the workflow and evaluate the test set

```
param_final <- rf_tune_results %>% select_best(metric = "roc_auc")
rf_workflow <- rf_workflow %>% finalize_workflow(param_final)
```

```
# fit on the training set and evaluate on test set
rf_fit <- rf_workflow %>% last_fit(forrest_split)
# check performance on test set
test_performance <- rf_fit %>% collect_metrics()
# generate predictions from the test set
test_predictions <- rf_fit %>% collect_predictions()
# create confusion matrix
confusion_matrix <- test_predictions %>% conf_mat(truth = isforrest, estimate = .pred_class)
confusion_matrix <- as.data.frame(confusion_matrix["table"])
```


One way to visualize the performance of a classification model is to use a confusion matrix.

```
plotTable <- confusion_matrix %>%
  mutate(goodbad = ifelse(confusion_matrix$table.Prediction == confusion_matrix$table.Truth, "good", "bad")) %>%
  group_by(table.Truth) %>%
  mutate(prop = table.Freq/sum(table.Freq))

ggplot(data = plotTable, mapping = aes(x = table.Truth, y = table.Prediction, fill = goodbad, alpha = prop)) +
  geom_tile() +
  geom_text(aes(label = table.Freq), vjust = .5, fontface = "bold", alpha = 1, color = "white") +
  theme_minimal() +
  ylim(rev(levels(confusion_matrix$table.Prediction))) +
  theme_ft_rc() +
  labs(title = "Confusion matrix",
       subtitle = "How well can we predict the speakers of spoken lines?",
       caption = "Data: Master Thesis") +
  theme(panel.border = element_blank(),
        panel.grid.major = element_blank(),
        panel.grid.minor = element_blank()) +
  scale_fill_npg(alpha = 1) +
  scale_color_npg()
```

Confusion matrix

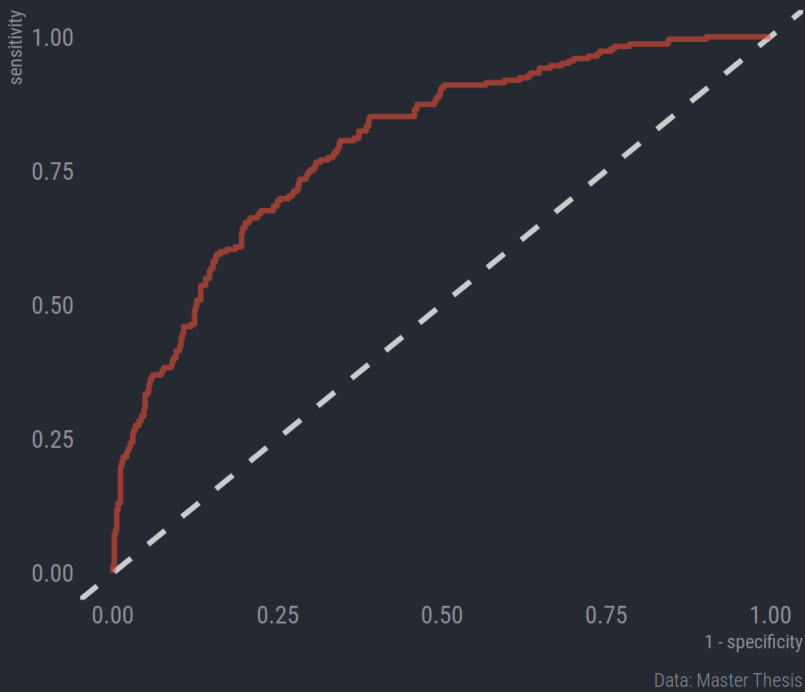
How well can we predict the speakers of spoken lines?



```
test_predictions %>%
  group_by(id) %>%
  roc_curve(isforrest, .pred_FORREST) %>%
  ggplot(aes(1 - specificity, sensitivity, color = id)) +
  geom_abline(lty = 2, color = "gray80", size = 1.2) +
  geom_path(show.legend = FALSE, alpha = 0.6, size = 1.2) +
  labs(title = "ROC curve",
       subtitle = "Performance of the selected random forest model",
       caption = "Data: Master Thesis") +
  theme_minimal() +
  theme_ft_rc() +
  theme(panel.border = element_blank(),
        panel.grid.major = element_blank(),
        panel.grid.minor = element_blank()) +
  scale_fill_npg() +
  scale_color_npg()
```

ROC curve

Performance of the selected random forest model



Our results here indicate that we did not overfit during the tuning process. We can also create a ROC curve for the testing set.

4 Extract important features

```
library(vip)
imp_df <- rf_workflow %>%
  fit(data = forrest_test) %>%
  pull_workflow_fit() %>%
  vi()

# rename variables for plotting
imp_df$Variable <- c("lexical frequency",
                    "lexical maximum",
                    "lexical minimum",
                    "speech rate",
                    "seconds",
                    "characters per word",
                    "lowercase characters",
                    "open words",
                    "unique_characters",
                    "closed words",
                    "syntactic nodes",
                    "sentiment",
                    "lexical diversity",
                    "uppercase characters",
                    "unique words",
                    "words",
                    "syntactic depth")

imp_df$Variable <- as.factor(imp_df$Variable)
```

Importance of predictors

These are the predictors that are most important globally for whether a line was spoken by Forrest or not.

